

Your AI agent can send that email. Should it?

A PM's framework for deciding when agents act alone
— and when humans stay in the loop.

MCP standardizes how agents use tools.

A2A standardizes how agents talk to each other.

Nobody standardizes when an agent should stop and ask a human.

That's not a protocol problem. It's a product decision.

Reversibility × Blast radius = Oversight level

Two questions decide everything: Can you undo it? Who does it reach?

REVERSIBLE + INTERNAL

Let the agent run. Interrupting the user here is pure friction.

REVERSIBLE + EXTERNAL

Act, then surface it. Humans review, they don't gate.

IRREVERSIBLE + INTERNAL

Pause for approval. Undo doesn't exist here.

IRREVERSIBLE + EXTERNAL

Highest stakes. Gate it — or design it out of the agent entirely.

T0

Go

Full autonomy. No human involvement.

Support agent searches the knowledge base to draft an answer

Meeting agent summarizes a call transcript

Analytics agent reads dashboards to compile a weekly digest

RULE Read-only and reversible? Never interrupt the user.

T1

Do, Then Tell

Act freely. Report after.

Coding agent opens a pull request — merge still needs review

CRM agent logs call notes and updates deal stages

Support agent tags and routes tickets, flags a daily summary

RULE Low-stakes writes, easily undone? Don't make humans a bottleneck.

T2

Ask, Then Act

Human approves before execution.

Email agent drafts the customer reply — human clicks send

Procurement agent fills the order — human confirms the purchase

Marketing agent schedules the post — human approves before it goes live

RULE Irreversible, external-facing, or spends money? Pause first.

T3

Hands Off

The agent never touches it. Human only.

Entering credentials or payment details

Signing contracts or accepting legal terms

Deleting production data or changing access permissions

RULE No approval flow makes these safe. Design them out entirely.

Classify any agent action in 10 seconds.

Check in order. First match wins.

Touches credentials, permissions, or legal?

T3 – agent never does it

Moves money, irreversible, or leaves the org?

T2 – ask, then act

Writes or changes anything — undoable, internal?

T1 – do, then tell

Read-only?

T0 – go

Design your agent to **stop, not guess.**

Approval times out? Instruction is ambiguous? Request denied twice? Spec it up front: the agent escalates to a human — it never acts on silence.



**Grant autonomy per action —
never per agent.**

MORE ON AI TOOLS FOR PM:

github.com/prashibadkur11-creator